

---

# Storage - Indexing

## Learning Maps

Sequence	Title
1	Introduction to databases
2	Relational Databases
3	Relational Algebra
4	Structured Query Language – Part 1
5	Structured Query Language – Part 2
6	Constraints and Triggers
7	Entity Relationship Model
8	Functional Dependency
9	Normalization
10	Storage - Indexing
11	Query Processing
12	Transaction Management – Part 1
13	Transaction Management – Part 2

## Intro > Overview



- ☐ A : Voice and PPT Overview
- ☐ B : Text-based Overview
- ☒ C : Video and PPT Overview

<b>Opening Message</b>	<p>→ In this lesson, we will study the physical data organization inside database management systems and the indexing mechanism to speed up SELECT queries.</p> <p>→ Storage and Indexing</p>
<b>Lesson topic</b>	<ol style="list-style-type: none"><li>1. Overview of database storage structures</li><li>2. Physical database file structures</li><li>3. Database index</li></ol>
<b>Learning Goals</b>	<p>Upon completion of this lesson, students will be able to:</p> <ol style="list-style-type: none"><li>1. Understand the physical database file structures</li><li>2. Understand the role of database indexes</li></ol>

## Intro > Keywords

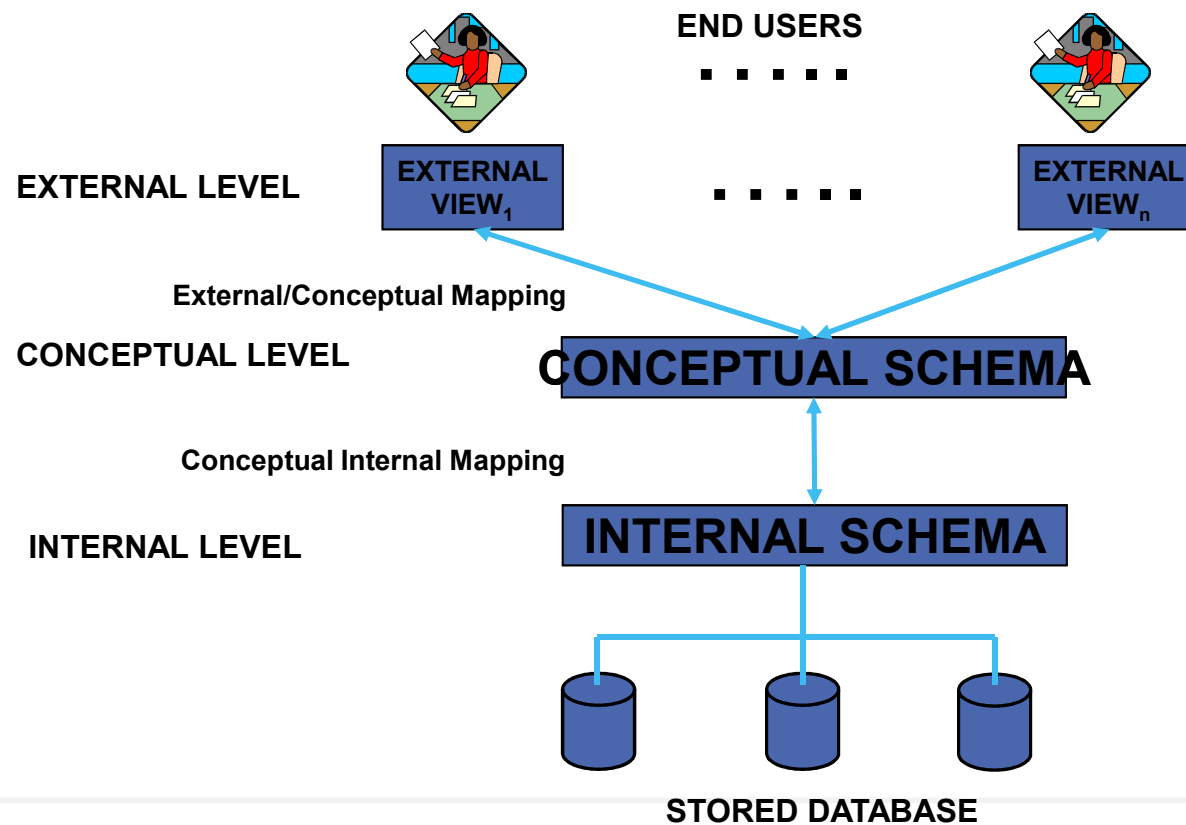
Keyword	Description
Heap file	Files of Unordered Records
Ordered file	Physically order the records of a file on disk based on the values of one of their fields (key field)
Index	A data structure that improves the speed of data retrieval operations
B-tree	A self-balancing tree data structure that keeps data sorted

## Lesson > Topic 1: Overview



- 3-tier Schema Model (ANSI-SPARC Architecture)
- How Mariadb stores data

## 1.1. 3-tier Schema Model (ANSI-SPARC Architecture)



## 1.2. How does Mariadb store data

```
MariaDB [(none)]> SHOW VARIABLES LIKE 'datadir';
```

Variable_name	Value
datadir	/var/lib/mysql/

```
MariaDB [student_management]> show tables;
```

Tables_in_student_management
class
enrolled
faculty
student

the .frm table file stores the table's format  
the .ibd file stores the table's data

```
:/var/lib/mysql/student_management# ls -la
```

ql mysql	4096	Mar 12 02:05	.
ql mysql	4096	May 5 06:06	
ql mysql	1547	Mar 12 02:05	class.frm
ql mysql	114688	Mar 12 02:21	class.ibd
ql mysql	65	Mar 12 01:59	db.opt
ql mysql	1466	Mar 12 02:03	enrolled.frm
ql mysql	114688	Mar 12 02:18	enrolled.ibd
ql mysql	1005	Mar 12 02:04	faculty.frm
ql mysql	98304	Mar 12 02:16	faculty.ibd
ql mysql	1101	Mar 12 02:00	student.frm
ql mysql	98304	Mar 12 02:23	student.ibd

## 1.2.1. the .frm file stores the table's format

```
MariaDB [student_management]> describe student;
```

Field	Type	Null	Key	Default	Extra
snum	int(11)	NO	PRI	NULL	
sname	varchar(40)	YES		NULL	
major	varchar(30)	YES		NULL	
level	varchar(10)	YES		NULL	
age	int(11)	YES		NULL	

```
root@285e07e9458f:/var/lib/mysql/student_management# cat student.frm
?
```

```
VM?\!  ?s?$??%?峰B??
??PRIMARY??InnoDB??f\P
          (/?
```

```
N?
```

```
?snum?sname?major?level?age?root@285e07e9458f:/var/lib/mysql/student
```



### 1.2.2. the .ibd file stores the table's data

```
MariaDB [student_management]> select * from student;
```

snum	sname	major	level	age
1	Nguyen Van A	CS	JR	18
2	Nguyen Viet Cuong	History	JR	19
3	Nguyen Hong Ngoc	CS	JR	19
4	Mark Juke	History	JR	20
5	Elon Mulk	CS	JR	20
6	Donal Trump	CS	JR	20
7	Obama	CS	JR	20
8	Tan Dung	History	SR	30

```
root@285e07e9458f:/var/lib/mysql/student_management# cat student.ibd
```

[illegible]

```

supremum
ong NgocCSJR? (0?1?1Mark JukeHistoryJR? 0+?0??E1on MuIKCSJR?
pCSJR?@'W?0?ObamaCSJRH????Tan DungHistorySR?pc??Q?'??root@285e07e9458f:/var/lib/mys

```

## Lesson > Topic 2: Physical file structures



- Motivation
- Magnetic disks as data storage
- Primary file organizations

## 2.1. Motivation

---

- Databases typically store large amounts of data **persistently** on disks
  - Databases are too large to fit entirely in main memory.
  - Disk - **nonvolatile storage** vs. Main memory - **volatile storage**
  - The cost of storage per unit is much cheaper

## 2.2. Store data on disks (magnetic disks)

---



- A disk is a *random access* addressable device.
- Transfer of data between main memory and disk takes place in units of disk blocks
- Typical disk block sizes: 4KB – 8KB
- Disk I/O (read/write from disk to main memory) overhead is the key factor of database performance optimization

## 2.3. Physical database file structures

---

- The process of **physical database design** involves choosing the particular data organization techniques that best suit the given application requirements (on SELECT, INSERT, UPDATE, DELETE)
- The data stored on disk is organized as **files of records**
  - **Primary file organizations:** determine how the file records are *physically placed* on the disk, *and hence how the records can be accessed*
  - **Secondary organization or auxiliary access structure** allows efficient access to file records based on *alternate fields*

## 2.4. Placing File Records on Disk



**Figure 17.5**

Three record storage formats. (a) A fixed-length record with six fields and size of 71 bytes. (b) A record with two variable-length fields and three fixed-length fields. (c) A variable-field record with three types of separator characters.

## 2.5. Primary file organizations

---

- Files of Unordered Records (Heap Files)
- Files of Ordered Records (Sorted Files)
- Hashing Techniques

## 2.5.1. Files of Unordered Records (Heap Files)

---

- Records are placed in the file in the order in which they are inserted
- INSERT: Inserting a new record is *very efficient*
  - New records are inserted at the end of the file
- UPDATE/SELECT: Searching for a record on any search condition is not efficient – linear search
- DELETE: leaves unused space in the disk block
  - require periodic **reorganization**



## 2.5.2. Files of Ordered Records (Sorted Files)

---

- Physically order the records of a file on disk based on the values of one of their fields (key field)
- SELECT: binary search (very fast)
- INSERT/DELETE/UPDATE: more expensive

### 2.5.3. Hash files

---

- The *address* of the disk block in which the record is stored is the result of applying a **hash function** to the value of a particular field (hash field) of the record
- Very fast access to records for search on equality condition on the **hash field**


## Lesson > Topic 3: Database indexes



- What is database indexes?
- Index data structures
- B+tree
- Sparse vs. Dense index
- Clustered vs. Non-clustered index
- Index creation in SQL

## 3.1. Secondary organization

- **Auxiliary access structure (commonly index)** allows efficient access to file records based on *alternate fields*




## 3.2. Index data structure

---

- Indexes can be implemented with different data structures.
  - **B+-tree index**
  - **hash index**
  - **bitmap index (briefly)**
  - dynamic hash indexes: number of buckets modified dynamically
  - R-tree: index for spacial data (points, lines, shapes)
  - quadtree: recursively partition a 2D plane into four quadrants
  - octree: quadtree version for three dimensional data
  - main memory indexes: T-tree, binary search tree

### 3.3. B+-Tree

- balanced tree of key-pointer pairs
- keys are sorted by value
- nodes are at least half full
- access records for key: traverse tree from root to leaf




## 3.4. Spare vs. Dense index

- **Sparse index**

- pointers to disk pages  
at most one pointer per disk page  
usually much less pointers than records


- **Dense index**

- pointers to individual records  
one key per record  
usually more keys than sparse index  
optimization: store repeating keys only once  
followed by pointers



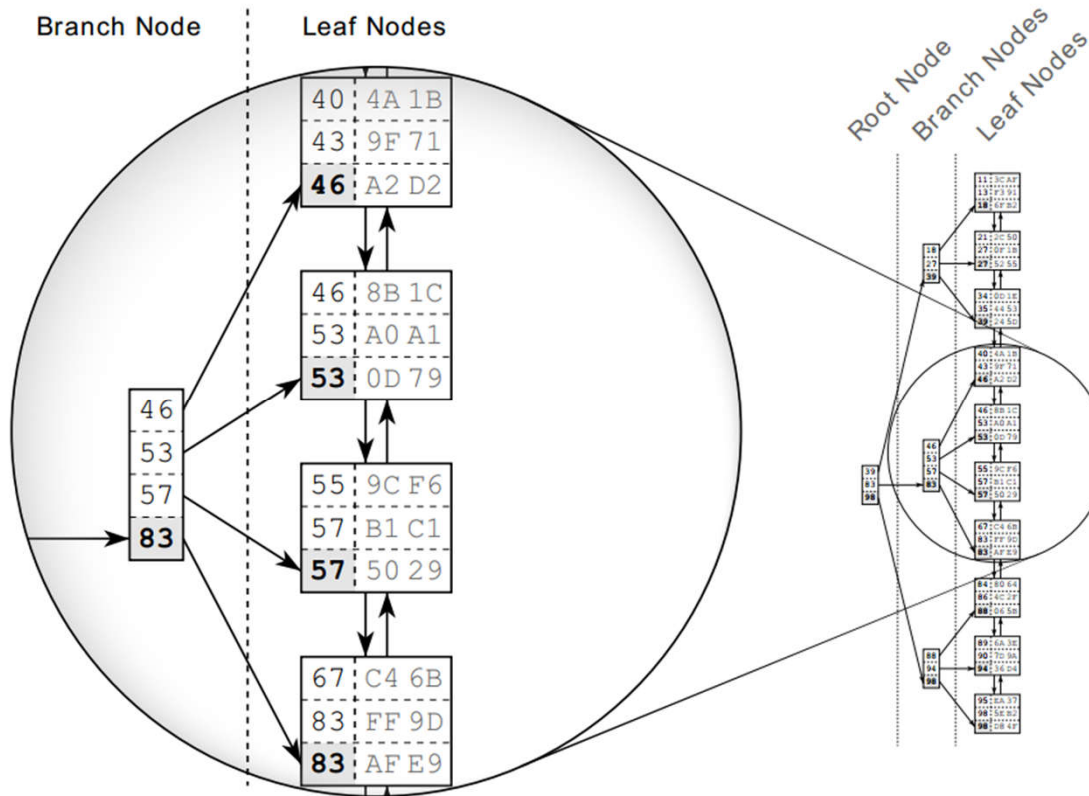
## 3.5. Clustered vs. Non-Clustered

- **Clustered index on attribute X**
  - This index controls the placement of records on disk
  - only one clustering index per table
  - dense or sparse
- **Non-clustered index on attribute X**
  - no constraint on table organization
  - Can have more than one index per table
  - always dense



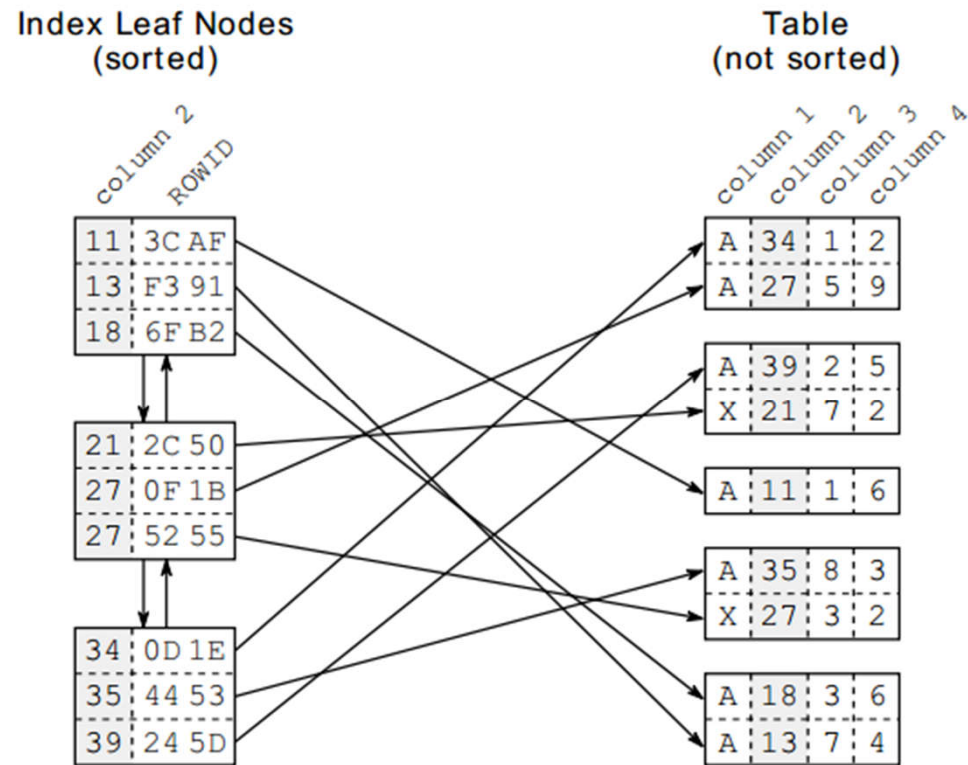


## 3.6. Example: B+ tree



© Gulutzan, Peter, and Trudy Pelzer. *SQL Performance Tuning*. Addison-Wesley Professional, 2003.

## 3.7. Example: Non-clustered index



## 3.8. CREATING INDEX

---

- CREATE [UNIQUE|FULLTEXT|SPATIAL] INDEX  
*index\_name* [*index\_type*] ON *tbl\_name*  
(*index\_col\_name*,...) [*index\_option*] [*algorithm\_option* |  
*lock\_option*] ...
- *index\_type*: USING {BTREE | HASH}

## Quiz



No	Question (Multiple Choice)	Answer (1,2,3,4)	Commentary
1	Does heap files support INSERT query efficiently? A. Yes B. No	A	
2	Are ordered files better for heavy Insert operation? A. Yes B. No	B	
3			

## Outro > Summary



No	Topic	Summary
1	Overview of database storage structures	3-tier Schema Model (ANSI-SPARC Architecture) How Mariadb stores data
2	Physical database file structures	Motivation Magnetic disks as data storage Primary file organizations
3	Database index	What is database indexes? Index data structures B+tree Spare vs. Dense index Clustered vs. Non-clustered index Index creation in SQL

---

You've just learned about storage and indexing

Next lesson:

## Query processing

1. Overview of query processing
  2. Query optimization
-