

---

# Query processing

## Learning Maps

Sequence	Title
1	Introduction to databases
2	Relational Databases
3	Relational Algebra
4	Structured Query Language – Part 1
5	Structured Query Language – Part 2
6	Constraints and Triggers
7	Entity Relationship Model
8	Functional Dependency
9	Normalization
10	Storage - Indexing
11	Query Processing
12	Transaction Management – Part 1
13	Transaction Management – Part 2

## Intro > Overview



- ☐ A : Voice and PPT Overview
- ☐ B : Text-based Overview
- ☒ C : Video and PPT Overview

<b>Opening Message</b>	<ul style="list-style-type: none"><li>→ In this lesson, we will study</li><li>→ Query processing</li><li>→ Query optimization</li></ul>
<b>Lesson topic</b>	<ul style="list-style-type: none"><li>1. Course overview</li><li>2. Basic concepts on database</li><li>3. Data management: database vs. file approach</li></ul>
<b>Learning Goals</b>	<p>Upon completion of this lesson, students will be able to:</p> <ul style="list-style-type: none"><li>1. Recall the concepts of database, DBMS, data model, file system.</li><li>2. Identify the characteristics of database and file system approach in data management</li></ul>

## Intro > Keywords

Keyword	Description
Query processing	Activities involved in retrieving/storing data from/to the database
Query optimization	Selection of an efficient query execution plan

## Lesson > Topic 1: Course overview



- Overview of query processing
- Query optimization

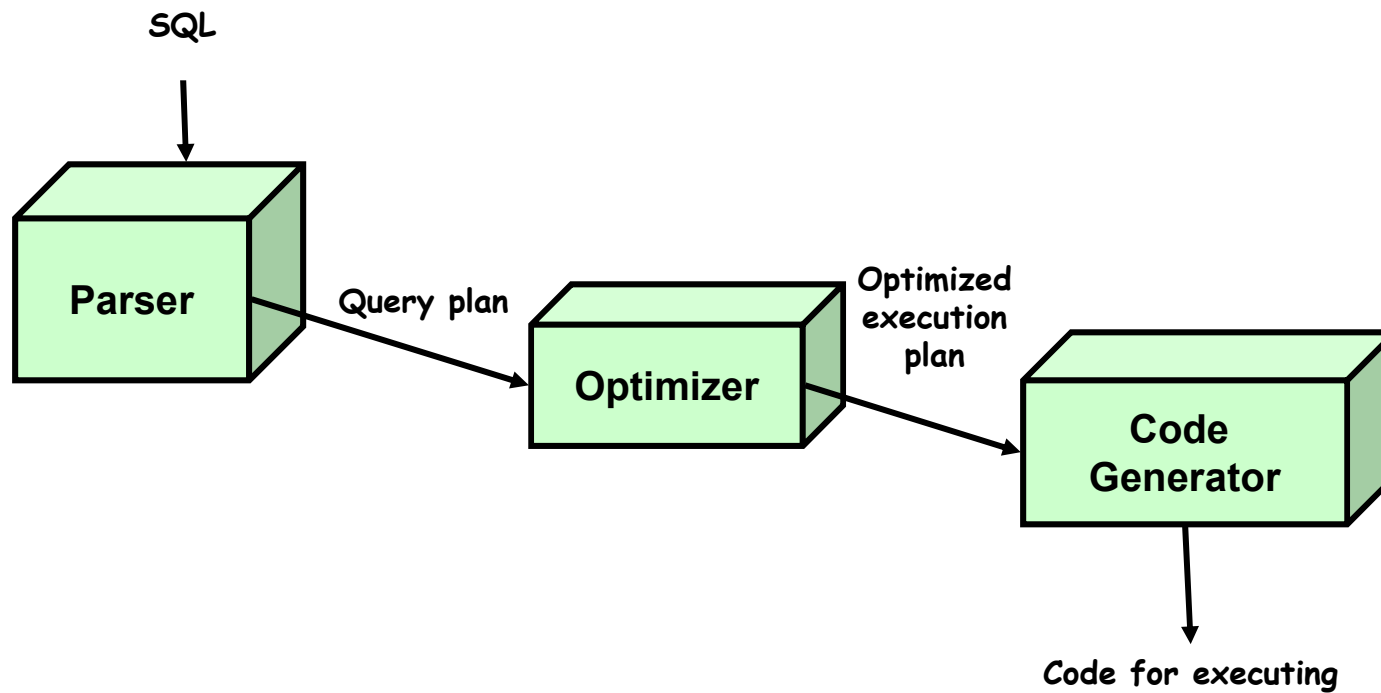
## 1.1. What is query processing

---

- The entire process or activities involved in retrieving data from the database
  - SQL query translation into low level instructions (usually relational algebra)
  - Query optimization to save resources, cost estimation or evaluation of query,
  - Query execution for the extraction of data from the database.

## 1.2. Phases of query processing

---



## 1.3. Parser

---

- Scans and parses the query into individual tokens and examines for the correctness of query
  - Does it contain the right keywords?
  - Does it conform to the syntax?
  - Does it contain the valid tables, attributes?
- Output: Query plan (Relational algebra)
- Input:
  - SELECT balance FROM account  
WHERE balance < 2500

- Output: Relational algebra (RA) expression

- But it's not unique

$$\Pi_{balance}(\sigma_{balance < 2500}(account))$$

$$\sigma_{balance < 2500}(\Pi_{balance}(account))$$




## 1.4. Optimizer

- Input: RA expression

$$\Pi_{balance}(\sigma_{balance < 2500}(account))$$

- Output: Query execution plan
  - Query execution plan = query plan + the algorithms for the executions of RA operations
  - Choose the cheapest execution plan out of the possible ones
    - Step 1: Equivalence transformation
    - Step 2: Annotation for the algorithm of the RA expression
    - Step 3: Cost estimation for different query execution plans



## 2. Understanding optimizer

- Choose the cheapest execution plan out of the possible ones
  - Step 1: Equivalence transformation
  - Step 2: Annotation for the algorithm of the RA expression
  - Step 3: Cost estimation for different query execution plans

## 2.1. Step 1: Equivalence transformation

---

- RA expressions are equivalent if they generate the same set of tuples on every database instance
- Equivalence rules:
  - Transform one relational algebra expression into equivalent one
  - Similar to numeric algebra:  $a + b = b + a$ ,  $a(b + c) = ab + ac$ , etc.
- Why producing equivalent expressions?
  - equivalent algebraic expressions give the same result
  - but usually the execution time varies significantly

## Equivalence transformation rules

---

- (1) Conjunctive selection operations can be deconstructed into a sequence of individual selections; cascade of  $\sigma$ 
  - $\sigma_{\theta_1 \wedge \theta_2}(E) = \sigma_{\theta_1}(\sigma_{\theta_2}(E))$
- (2) Selection operations are commutative
  - $\sigma_{\theta_1}(\sigma_{\theta_2}(E)) = \sigma_{\theta_2}(\sigma_{\theta_1}(E))$
- (3) Only the final operations in a sequence of projection operations is needed; cascade of  $\Pi$ 
  - $\Pi_{L_1}(\Pi_{L_2}(\dots \Pi_{L_n}(E) \dots)) = \Pi_{L_1}(E)$
- (4) Selections can be combined with Cartesian products and theta joins
  - $\sigma_{\theta_1}(E_1 \times E_2) = E_1 \bowtie_{\theta_1} E_2$
  - $\sigma_{\theta_1}(E_1 \bowtie_{\theta_2} E_2) = E_1 \bowtie_{\theta_1 \wedge \theta_2} E_2$

- 
- (5) Theta Join operations are commutative
    - $E_1 \bowtie_{\theta} E_2 = E_2 \bowtie_{\theta} E_1$
  - (6) Natural join operations are associative
    - $E_1 \bowtie (E_2 \bowtie E_3) = (E_1 \bowtie E_2) \bowtie E_3$
    - Theta join are associative in the following manner where  $\theta_2$  involves attributes from  $E_2$  and  $E_3$  only
    - $(E_1 \bowtie_{\theta_1} E_2) \bowtie_{\theta_2 \wedge \theta_3} E_3 = E_1 \bowtie_{\theta_1 \wedge \theta_3} (E_2 \bowtie_{\theta_2} E_3)$

- 
- (7) Selection distributes over joins in the following ways
    - If predicate involves attributes of  $E_1$  only
    - $\sigma_{\theta_1} (E_1 \bowtie_{\theta_2} E_2) = \sigma_{\theta_1} (E_1) \bowtie_{\theta_2} E_2$
    - If predicate  $\theta_1$  involves only attributes of  $E_1$  and  $\theta_2$  involves only attributes of  $E_2$  (a consequence of rule 7 and 1)
    - $\sigma_{\theta_1 \wedge \theta_2} (E_1 \bowtie_{\theta_3} E_2) = \sigma_{\theta_1} (E_1) \bowtie_{\theta_3} \sigma_{\theta_2} (E_2)$

- 
- (8) Projection distributes over join as follows
    - $\Pi_{L_1 \cup L_2}(E_1 \bowtie_{\theta} E_2) = \Pi_{L_1}(E_1) \bowtie_{\theta} \Pi_{L_2}(E_2)$
    - If  $\theta$  involves attributes in  $L_1 \cup L_2$  only and  $L_i$  contains attributes of  $E_i$
  - (9) The set operations union and intersection are commutative
    - $E_1 \cup E_2 = E_2 \cup E_1$
    - $E_1 \cap E_2 = E_2 \cap E_1$
  - (10) The union and intersection are associative
    - $(E_1 \cup E_2) \cup E_3 = E_1 \cup (E_2 \cup E_3)$

- 
- (11) The selection operation distributes over union, intersection, and set-difference
    - $\sigma_{\theta}(E_1 \cup E_2) = \sigma_{\theta}(E_1) \cup \sigma_{\theta}(E_2)$
    - $\sigma_{\theta}(E_1 \cap E_2) = \sigma_{\theta}(E_1) \cap \sigma_{\theta}(E_2)$
    - $\sigma_{\theta}(E_1 - E_2) = \sigma_{\theta}(E_1) - \sigma_{\theta}(E_2)$
  - (12) The project operation distributes over the union
    - $\Pi_L(E_1 \cup E_2) = \Pi_L(E_1) \cup \Pi_L(E_2)$



## 2.2. Step 2: Execution algorithms of RA operations

---

- Algebra expression is not a query execution plan.
- Additional decisions required:
  - which indexes to use, for example, for joins and selects?
  - which algorithms to use, for example, sort-merge vs. hash join?
  - materialize intermediate results or pipeline them?

## 2.2.1. Basic Operators

---

- One-pass operators:
  - Scan
  - Select
  - Project
- Multi-pass operators:
  - Join
    - Various implementations
    - Handling of larger-than-memory sources
  - Semi-join
  - Aggregation, union, etc.

## 2.2.2. 1-Pass Operators: Scanning a Table

---

- Sequential scan: read through blocks of table
- Index scan: retrieve tuples in index order

### 2.2.3. Nested-loop JOIN

---

```
for each tuple tr in r {  
  for each tuple ts in s {  
    if (tr and ts satisfy the join condition  $\theta$ ) {  
      add tuple  $tr \times ts$  to the result set  
    }  
  }  
}
```

- No index needed
- Any join condition types
- Expensive:  $O(n^2)$

## 2.2.4. Single-loop JOIN (Index-based)

---

```
for each tuple tr in R {  
  ts = index.get(tuple s) {  
    if ts.exist() {  
      add tr x ts to the result set  
    }  
  }  
}
```

- Only if there is at least one index on join attributes


## 2.2.5. Sort-merge JOIN

---

- Requires data physically sorted by join attributes
  - Merge and join sorted files, reading sequentially a block at a time
    - Maintain two file pointers
      - While tuple at R < tuple at S, advance R (and vice versa)
      - While tuples match, output all possible pairings
    - Preserves sorted order of “outer” relation
- ✓ Very efficient for presorted data
- ✓ Can be “hybridized” with NL Join for range joins
- × May require a sort before (adds cost + delay)
- Cost:  $b(R) + b(S)$  plus sort costs, if necessary  
In practice, approximately linear,  $3 (b(R) + b(S))$

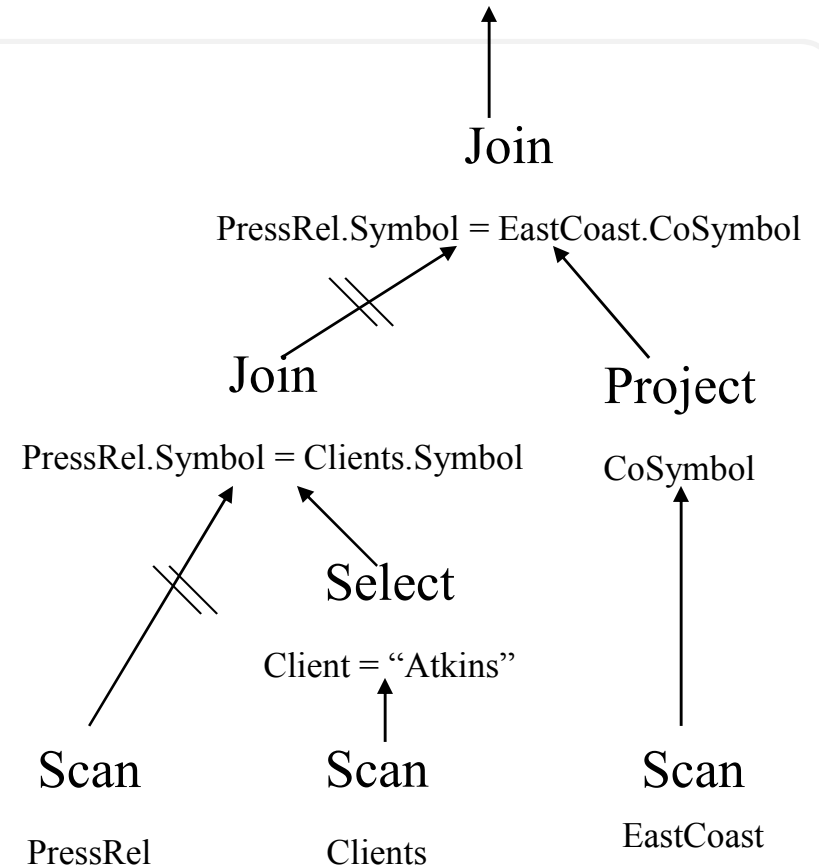
## 2.2.6. Partition-hash JOIN

- Hash two relations on join attributes
- Join buckets accordingly



## 2.2.7. Execution Strategy: Materialization vs. Pipelining

- Execution strategy defines how to walk the query execution plan
  - Materialization
  - Pipelining





## Materialization

---

- Performs the innermost or leaf-level operations first of the query execution plan
- The intermediate result of each operation is materialized into temporary relation and becomes input for subsequent operations.
- The cost of materialization is the sum of the individual operations plus the cost of writing the intermediate results to disk
  - lots of temporary files, lots of I/O.

# Pipelining

---

- Operations form a queue, and results are passed from one operation to another as they are calculated
- Pipelining restructures the individual operation algorithms so that they take streams of tuples as both input and output.
- Limitation
  - algorithms that require sorting can only use pipelining if the input is already sorted beforehand
  - since sorting by nature cannot be performed until all tuples to be sorted are known.

## 2.3. Step 3: Cost estimation

---

- Each relational algebra expression can result in many query execution plans
- Some query execution plans may be better than others
- Finding the fastest one
  - Just an estimation under certain assumptions
  - Huge number of query plans may exist

## Cost estimation factors

---

- Catalog information: database maintains statistics about relations
- Ex.
  - number of tuples per relation
  - number of blocks on disk per relation
  - number of distinct values per attribute
  - histogram of values per attribute
- Problems
  - cost can only be estimated
  - updating statistics is expensive, thus they are often out of date

## Choosing the cheapest query plan

---

- Problem: Estimating cost for all possible plans too expensive.
- Solutions:
  - pruning: stop early to evaluate a plan
  - heuristics: do not evaluate all plans
- Real databases use a combination of
  - Apply heuristics to choose promising query plans.
  - Choose cheapest plan among the promising plans using pruning.
- Examples of heuristics:
  - perform selections as early as possible
  - perform projections early avoid Cartesian products

# Heuristic rules

- ▶ Query optimizers use the equivalence rules of relational algebra to improve the expected performance of a given query in *most cases*.
- ▶ The optimization is guided by the following **heuristics**:
  - (a) **Break apart conjunctive selections** into a sequence of simpler selections  
(rule ①—preparatory step for (b)).
  - (b) **Move  $\sigma$  down the query tree** for the earliest possible execution  
(rules ②, ⑦, ⑪—reduce number of tuples processed).
  - (c) **Replace  $\sigma$ - $\times$  pairs by  $\bowtie$**   
(rule ④ (a)—avoid large intermediate results).
  - (d) **Break apart and move as far down the tree as possible lists of projection attributes**, create new projections where possible  
(rules ③, ⑧, ⑫—reduce tuple widths early).
  - (e) **Perform the joins with the smallest expected result first.**

## Quiz



No	Question (Multiple Choice)	Answer (1,2,3,4)	Commentary
1	<p>The reaction <math>A \rightarrow P</math> follow the first-order rule. The half-life of the reaction is 60 minutes. The time for A concentration of 12.5 %</p> <p>A. 60 min B. 30 min C. 120 min D. 180 min</p>	D	3 times of half-life
2	<p>Consider the combustion of propene:  <math>2C_3H_6 + 9O_2 \rightarrow 6CO_2 + 6H_2O</math>.                      If the formation rate of <math>CO_2</math> is <math>0.30 \text{ M.s}^{-1}</math>, the consumption rate of <math>C_3H_6</math> is:</p> <p>A. <math>0.90 \text{ M.s}^{-1}</math>                      B. <math>0.10 \text{ M.s}^{-1}</math>                      C. <math>0.45 \text{ M.s}^{-1}</math>                      D. <math>0.30 \text{ M.s}^{-1}</math></p>	A	<p>Reaction rate:  <math>\frac{1}{6} \frac{d[CO_2]}{dt} = -\frac{1}{2} \frac{d[C_3H_6]}{dt}</math>                      Formation rate of <math>CO_2</math>:  <math>R_{CO_2} = \frac{d[CO_2]}{dt}</math>                      Consumption rate of <math>C_3H_6</math>:  <math>R_{C_3H_6} = -\frac{d[C_3H_6]}{dt}</math></p>
3			

## Outro > Summary



No	Topic	Summary
1	Introduction to Chemical Kinetics	- While the negative change in Gibbs energy (or Helmholtz free energy)
2		
3		
4		